

Linear regression

Simple linear regression


Code:

```
import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
from sklearn import linear_model
from sklearn.metrics import r2_score

df=pd.read_csv('Fuel.csv')
df.head()

df.describe()# count, mean,min,25%,50%,75%
values = df[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_COMB','CO2EMISSIONS']]

# plot the relationship between engine size and co2 emission
plt.figure(figsize=(12,7))
plt.scatter(values['ENGINE SIZE'],values['CO2EMISSIONS'])
plt.xlabel('engine size')
plt.ylabel('co2')
```



```
#plot the relationship between cylinder and co2 emission
```

```
plt.figure(figsize=(12,7))
```

```
plt.scatter(values['CYLINDERS'],values['CO2EMISSIONS'])
```

```
plt.xlabel('cylinder size')
```

```
plt.ylabel('co2')
```

```
# get train and test data, train: 80%, test: 20%
```

```
msk=np.random.rand(len(values))<0.8
```

```
train=values[msk]
```

```
test=values[~msk]# ~ works for array not list
```

```
# linear regression model
```

```
linear=linear_model.LinearRegression()
```

```
train_x=np.array(train['ENGINE SIZE']).reshape(-1,1)
```

```
train_y=np.array(train['CO2EMISSIONS']).reshape(-1,1)
```

```
linear.fit(train_x,train_y) # fit(x,y)
```

```
print('the coef is {0}, the intercept is {1}'.format(linear.coef_,linear.intercept_))
```

```
# plot predicted data
```


```
plt.figure(figsize=(12,7))
```

```
plt.scatter(train['ENGINE SIZE'],train['CO2EMISSIONS'])
```

```
plt.plot(train_x,linear.coef_[0][0]*train_x + linear.intercept_[0],color='r')
```

```
plt.xlabel('engine')
```

```
plt.ylabel('Emission')
```



```
# test the value
```

```
test_x=np.array(test['ENGINE SIZE']).reshape(-1,1)
```

```
test_y=np.array(test['CO2EMISSIONS']).reshape(-1,1)
```

```
predict_y=linear.predict(test_x)
```

```
print('The mean of square error is {}'.format(np.mean((predict_y-test_y)**2)))
```

```
print('R2 is {}'.format(r2_score(test_y,predict_y))) # R2 the bigger the better
```

Multiple linear regression

Code:

```
# import modules
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from sklearn import linear_model

data=pd.read_csv('Fuel.csv')
data.head()

cdata=data[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_CITY','FUELCONSUMPTION_HWY','FUELCONSUMPTION_COMB','CO2EMISSIONS']]


# generate training and testing data: 80% for train
msk=np.random.rand(len(data))<0.8
train=cdata[msk]
test=cdata[~msk]

# use linear regression model
regr=linear_model.LinearRegression()
x=np.asanyarray(train[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_COMB']])
y=np.asanyarray(train[['CO2EMISSIONS']])
regr.fit(x,y)
print('the coefficient is: {}'.format(regr.coef_))
```

```

plt.figure(figsize=(12,7))
y_pred=regr.predict(test[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_COMB']])
test_x=np.asanyarray(test[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_COMB']])
test_y=np.asanyarray(test[['CO2EMISSIONS']])
plt.plot(y_pred)
plt.plot(test_y,color='r')
plt.xlabel('different factors')
plt.ylabel('CO2 emission')

```



```

# evaluate the model
print('mean of error square is {}'.format(np.mean((y_pred-test_y)**2)))
print('score is {}'.format(regr.score(test_x,test_y)))

```

Nonlinear regression

code:

```
import numpy as np
```

```

import matplotlib.pyplot as plt
from scipy.optimize import curve_fit
import pandas as pd


df=pd.read_csv('china_gdp.csv')
plt.plot(df['Year'],df['Value'])
plt.xlabel('Years')
plt.ylabel('GDP')

x=df['Year'].values
y=df['Value'].values

# choose a function and get the parameters
# in this case, logistic function is a good choice, but nomalization is required
def f(x,b1,b2):
    y=1/(1+np.exp(-b1*(x-b2)))
    return y
x_norm=x/max(x)
y_norm=y/max(y)
popt,pcov=curve_fit(f,x_norm,y_norm)
print('the parameters are: ',popt)

# plot the result
plt.figure(figsize=(12,7))
plt.scatter(x_norm,y_norm,label='actual',color='b')
y_p=f(x_norm,*popt)
plt.plot(x_norm,y_p,label='fit',color='r')
plt.legend(loc='best')
plt.xlabel('Year')
plt.ylabel('GDP')

```



Classification

Knn

Code:

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Fri Dec 14 22:24:12 2018

@author: hejia

```
"""
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn import metrics

# import data
df=pd.read_csv('C:\\Users\\hejia\\Documents\\python\\machine learning\\teleCust1000t.csv')

all_columns=df.columns

# series method value_counts()
df['gender'].value_counts()

# use hist to see distribution
plt.hist(df['income'])

for column in all_columns:

    plt.figure(figsize=(10,6))

    plt.hist(df[column],label=column)

    plt.legend(loc='best')


# convert Pandas data frame to Numpy array
X=df[all_columns[0:-1]].values
y=df[all_columns[-1]].values


# Normalize data
from sklearn import preprocessing
X=preprocessing.StandardScaler().fit(X).transform(X.astype(float))


# train and test split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=4)


#train and predict
k=4

knn=KNeighborsClassifier(k)

neigh=knn.fit(X_train,y_train)

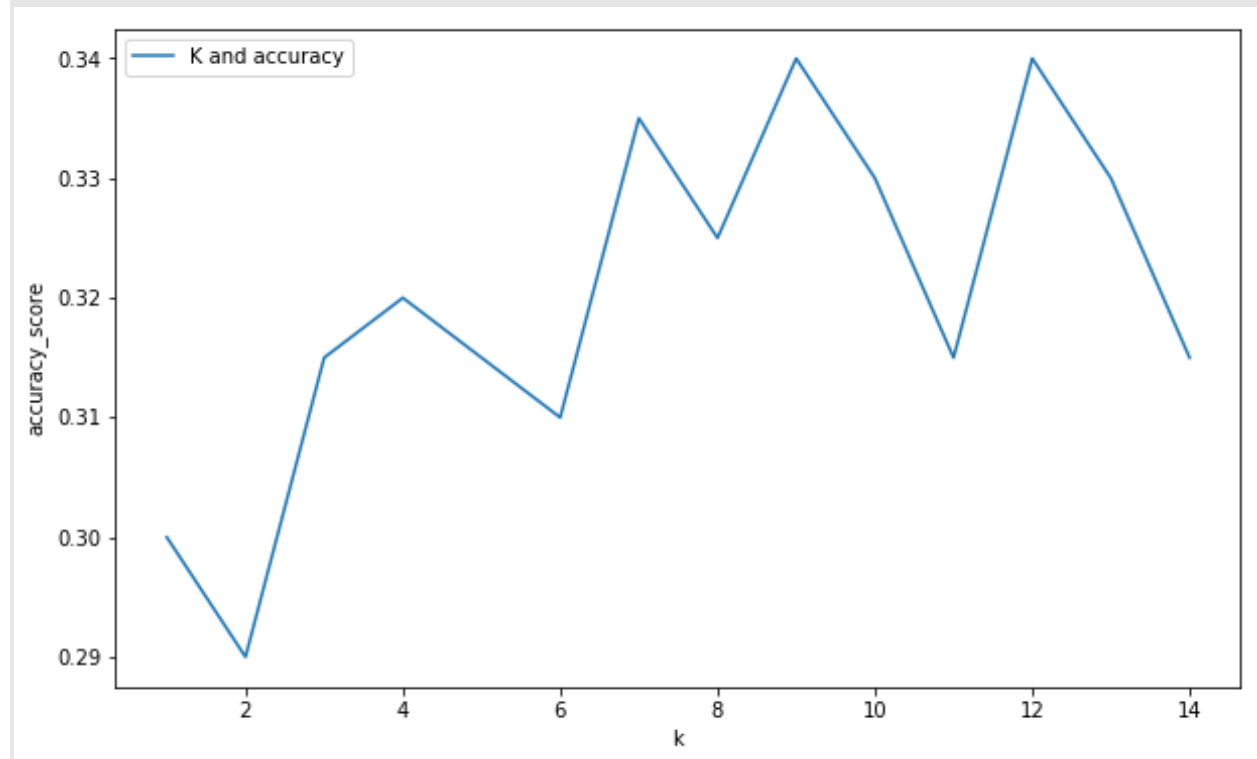
y_pred=neigh.predict(X_test)
```



```
# evaluation
accuracy=metrics.accuracy_score(y_test,y_pred)

# evaluate the relationship between k and accuracy
accuracy=[]
for k in range(1,15):
    knn=KNeighborsClassifier(k)
    y_pred=knn.fit(X_train,y_train).predict(X_test)
    accuracy.append(metrics.accuracy_score(y_test,y_pred))

plt.figure(figsize=(10,6))
plt.plot(range(1,15),accuracy,label='K and accuracy')
plt.xlabel('k')
plt.ylabel('accuracy_score')
plt.legend(loc='best')
```




Decision trees

Logic:

Decision tree learning algorithm

1. Choose an attribute from your dataset.
2. Calculate the significance of attribute in splitting of data.
3. Split data based on the value of the best attribute.
4. Go to step 1.



Code:

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Sun Dec 16 20:04:58 2018

@author: hejia

```
"""
```

```
import numpy as np
```

```
import pandas as pd
```

```
from sklearn import preprocessing
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn import metrics
```

```
# import data
```

```
df=pd.read_csv('drug200.csv')
```

```
columns=df.columns
X=df[columns[:-1]].values
y=df[columns[-1]]

# transfer text values to numerical
sex_code=preprocessing.LabelEncoder()
sex_code.fit(['F','M'])
X[:,1]=sex_code.transform(X[:,1])

BP_code=preprocessing.LabelEncoder()
BP_code.fit(['LOW','NORMAL','HIGH'])
X[:,2]=BP_code.transform(X[:,2])

chol_code=preprocessing.LabelEncoder()
chol_code.fit(['NORMAL','HIGH'])
X[:,3]=chol_code.transform(X[:,3])

X_trainset, X_testset, y_trainset, y_testset=train_test_split(X,y,test_size=0.3,random_state=3)

# model with decision tree
drugTree=DecisionTreeClassifier(criterion='entropy',max_depth=4)
drugTree.fit(X_trainset,y_trainset)
predTree=drugTree.predict(X_testset)

# evaluation
print('decision tree accuracy: ',metrics.accuracy_score(y_testset,predTree))
```

Logistic regression

Code:

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Mon Dec 17 05:09:50 2018

@author: hejia

```
"""
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import preprocessing
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
# import data
```

```
df=pd.read_csv('ChurnData.csv')
```

```
columns=['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip', 'callcard', 'wireless', 'churn']
```

```
df=df[columns]
```

```
df['churn']=df['churn'].astype('int') # convert the data type
```

```
X=df[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip']].values
```

```
y=df['churn'].values
```

```
#preprocess the data
```

```
X=preprocessing.StandardScaler().fit(X).transform(X)
```

```
X_train,X_test, y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=4)
```

```
logreg=LogisticRegression().fit(X_train,y_train)
```

```
y_pred=logreg.predict(X_test)
```

Cluster

K means

Code:

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Mon Dec 17 20:23:03 2018

@author: hejia

```
"""
```

```
import random
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.cluster import KMeans
```

```
from sklearn.datasets.samples_generator import make_blobs
```

```
# generate source data
```

```
np.random.seed(0)
```

```
X, y = make_blobs(n_samples=5000, centers=[[4,4], [-2, -1], [2, -3], [1, 1]], cluster_std=0.9)
```

```
plt.scatter(X[:,0],X[:,1],marker='.')
```

```
# set up K means
```

```
k_means=KMeans(n_clusters=4,n_init=12)
```

```
k_means.fit(X)
```

```
k_means_labels=k_means.labels_
```

```
k_means_labels
```

```


k_means_cluster_centers=k_means.cluster_centers_

# plot the data
fig=plt.figure(figsize=(8,5))
colors = plt.cm.Spectral(np.linspace(0, 1, len(set(k_means_labels))))
ax=fig.add_subplot(1,1,1)

for k,col in zip(range(0,4),colors):
    my_members=(k_means_labels==k)
    cluster_center=k_means_cluster_centers[k]
    # select for the points belonging to cluster K
    ax.plot(X[my_members, 0], X[my_members, 1], 'w', markerfacecolor=col, marker='.')
    ax.plot(cluster_center[0], cluster_center[1], 'o', markerfacecolor=col, markeredgecolor='k',
    markersize=6)

# Title of the plot
ax.set_title('KMeans')

```



```
# Remove x-axis ticks
ax.set_xticks(())

# Remove y-axis ticks
ax.set_yticks(())

# Show the plot
plt.show()

# use the real data
df=pd.read_csv('Cust_Segmentation.csv')

# address is not used
df.drop('Address',axis=1,inplace=True)

from sklearn.preprocessing import StandardScaler
X = df.values[:,1:]
X = np.nan_to_num(X) # replace nan with zero
Clus_dataSet = StandardScaler().fit_transform(X)
Clus_dataSet


# K mean modeling
clusterNum = 3
k_means = KMeans(init = "k-means++", n_clusters = clusterNum, n_init = 12)
k_means.fit(X)
labels = k_means.labels_
print(labels)

df['Clus_km']=labels
df.groupby('Clus_km').mean()
```

```

area = np.pi * ( X[:, 1])**2
plt.scatter(X[:, 0], X[:, 3], s=area, c=labels.astype(np.float), alpha=0.5)
plt.xlabel('Age', fontsize=18)
plt.ylabel('Income', fontsize=16)
plt.scatter()

```



```

from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure(1, figsize=(8, 6))
plt.clf()
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=48, azim=134)

plt.cla()
# plt.ylabel('Age', fontsize=18)
# plt.xlabel('Income', fontsize=16)
# plt.zlabel('Education', fontsize=16)


```



```

ax.set_xlabel('Education')
ax.set_ylabel('Age')
ax.set_zlabel('Income')
ax.scatter(X[:, 1], X[:, 0], X[:, 3], c= labels.astype(np.float))

```



Hierarchical

Code:

```

import numpy as np
import pandas as pd
from scipy import ndimage
from scipy.cluster import hierarchy
from scipy.spatial import distance_matrix
from matplotlib import pyplot as plt

```

```

from sklearn import manifold, datasets

from sklearn.cluster import AgglomerativeClustering


from sklearn.datasets.samples_generator import make_blobs

# generate data using make_blobs function

X1, y1 = make_blobs(n_samples=50, centers=[[4,4], [-2, -1], [1, 1], [10,4]], cluster_std=0.9)

plt.scatter(X1[:, 0], X1[:, 1], marker='o')

```



```

# choose how many clusters to form

agglom = AgglomerativeClustering(n_clusters = 4, linkage = 'average')

agglom.fit(X1,y1)

# Create a figure of size 6 inches by 4 inches.

plt.figure(figsize=(6,4))

# These two lines of code are used to scale the data points down,
# Or else the data points will be scattered very far apart.

```


```

# Create a minimum and maximum range of X1.
x_min, x_max = np.min(X1, axis=0), np.max(X1, axis=0)

# Get the average distance for X1.
X1 = (X1 - x_min) / (x_max - x_min)

# This loop displays all of the datapoints.
for i in range(X1.shape[0]):
    # Replace the data points with their respective cluster value
    # (ex. 0) and is color coded with a colormap (plt.cm.spectral)
    plt.text(X1[i, 0], X1[i, 1], str(y1[i]),
             color=plt.cm.nipy_spectral(agglom.labels_[i] / 10.),
             fontdict={'weight': 'bold', 'size': 9})

```



```

# Remove the x ticks, y ticks, x and y axis
plt.xticks([])
plt.yticks([])
#plt.axis('off')

# Display the plot of the original data before clustering

```

```
plt.scatter(X1[:, 0], X1[:, 1], marker='.')
```

